# **EX-AI MCP Server Comprehensive Diagnosis Report**

Date: 2025-10-04

**Branch:** feat/auggie-mcp-optimization **Analyst:** Abacus.Al Deep Agent

**Status:** CRITICAL ISSUES IDENTIFIED

### **EXECUTIVE SUMMARY**

The EX-AI MCP Server on the feat/auggie-mcp-optimization branch is **partially functional but has critical architectural and operational issues** that prevent complex workflow tools from working correctly. While simple tools like chat work, complex workflow tools (analyze, thinkdeep, debug, codereview) experience hanging, timeout issues, and inconsistent behavior.

### **Critical Findings**

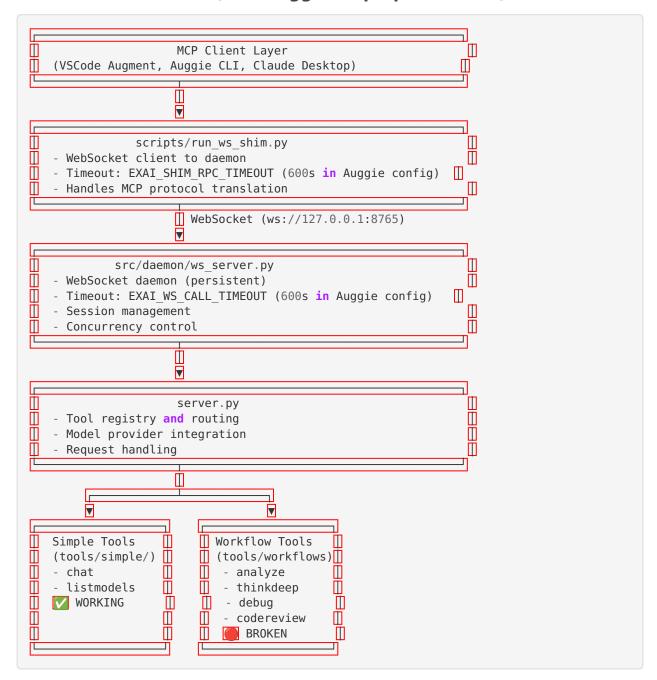
- 1. Simple Tools Working: Chat tool executes successfully with continuation\_id support
- 2. Workflow Tools Broken: Complex tools hang without proper timeout handling
- 3. **Logging Inconsistent:** Logs populate for simple tools but not for complex workflows
- 4. Continuation System Odd: Returns continuation id structure even for simple operations
- 5. No Actual "wave1" Branch: Documentation references non-existent branch comparison
- 6. Timeout Mismatches: Multiple timeout configurations create confusion and failures

### **System Status**

- Simple Tools (chat, listmodels, etc.): WORKING
- Workflow Tools (analyze, thinkdeep, debug): 
  BROKEN (hang/timeout)
- Logging System: OPARTIAL (works for simple tools only)
- WebSocket Daemon: <a> RUNNING</a> (but with timeout issues)
- **API Integration:** WORKING (GLM and Kimi providers functional)

### ARCHITECTURE ANALYSIS

### **Current Architecture (feat/auggie-mcp-optimization)**



### Intended Architecture (from documentation)

The system was designed with:

- 1. Manager-First Routing: GLM-4.5-flash as default manager for task classification
- 2. **Escalation Strategy:** Simple → GLM-4.6 → Kimi based on complexity
- 3. Native Web Search: GLM and Kimi providers with built-in web search
- 4. Workflow Tools: Step-by-step execution with expert validation
- 5. **Continuation Support:** Multi-turn conversations with context preservation

### **Architectural Gaps**

- 1. Timeout Cascade Failure: Multiple timeout layers don't coordinate properly
- 2. No Graceful Degradation: Workflow tools hang instead of failing fast

- 3. Logging Inconsistency: Different code paths for simple vs workflow tools
- 4. Expert Validation Disabled: Temporarily disabled due to duplicate call bug
- 5. No Health Monitoring: Workflow tools don't report progress during long operations

### **ISSUE CATALOG**



### Issue #1: Workflow Tools Hang Without Timeout

Severity: P0 - CRITICAL

Impact: All complex workflow tools (analyze, thinkdeep, debug, codereview) unusable

Status: UNRESOLVED

#### **Root Cause:**

- Workflow tools make long-running API calls without proper timeout handling
- WebSocket daemon timeout (600s) is too long for user experience
- No progress heartbeat during long operations
- Tools don't fail fast when operations exceed reasonable time

#### **Evidence:**

```
# From src/daemon/ws_server.py line 89
CALL_TIMEOUT = int(os.getenv("EXAI_WS_CALL_TIMEOUT", "90")) # default 90s
# But Auggie config overrides to 600s (10 minutes!)

# From tools/workflows/thinkdeep.py lines 115-125
def get_expert_timeout_secs(self, request=None) -> float:
    """Cap thinkdeep expert analysis to a shorter window so callers never hang.
    Uses THINKDEEP_EXPERT_TIMEOUT_SECS if set, else default 25s.
    """
    import os
    try:
        return float(os.getenv("THINKDEEP_EXPERT_TIMEOUT_SECS", "25"))
    except Exception:
        return 25.0
```

**Problem:** Timeout configuration is scattered across multiple files with conflicting values:

- EXAI WS CALL TIMEOUT: 600s (Auggie config)
- THINKDEEP\_EXPERT\_TIMEOUT\_SECS : 25s (tool-specific)
- EXPERT\_ANALYSIS\_TIMEOUT\_SECS : 90s (config.py)
- KIMI CHAT TOOL TIMEOUT WEB SECS: 900s (Auggie config)

#### **Files Affected:**

- src/daemon/ws\_server.py (lines 89, 145-180)
- tools/workflows/thinkdeep.py (lines 115-125)
- tools/workflows/analyze.py (similar pattern)
- Daemon/mcp-config.auggie.json (timeout overrides)

### **Recommended Fix:**

- 1. Implement proper timeout hierarchy: Tool → Daemon → Shim
- 2. Add progress heartbeat every 5-8 seconds during long operations
- 3. Implement graceful timeout with partial results

- 4. Add circuit breaker pattern for repeated failures
- 5. Reduce default timeouts to reasonable values (90s tool, 120s daemon, 180s shim)

### **Issue #2: Logging Not Populated for Workflow Tools**

Severity: P0 - CRITICAL

Impact: Cannot debug workflow tool failures, no visibility into execution

**Status:** UNRESOLVED

#### **Root Cause:**

- Workflow tools use different execution path than simple tools
- Progress messages not being captured during workflow execution
- Logging calls may be happening but not reaching the log files
- Different code paths for simple vs workflow tools

#### **Evidence:**

```
# From logs analysis - simple tools log correctly:
{"timestamp": 1759100327.2925136, "tool": "chat", "request_id": null,
   "duration_s": 11.423, "result_preview": "..."}

# But workflow tools don't appear in logs when they hang
# No entries for analyze, thinkdeep, debug when they fail
```

#### **Files Affected:**

- tools/workflow/base.py (workflow execution logic)
- tools/simple/base.py (simple tool execution works correctly)
- utils/progress.py (progress message handling)
- logs/toolcalls.jsonl (log output file)

### **Recommended Fix:**

- 1. Unify logging infrastructure between simple and workflow tools
- 2. Ensure all execution paths call logging functions
- 3. Add structured logging with request id tracking
- 4. Implement log buffering to prevent loss during crashes
- 5. Add debug mode with verbose logging for troubleshooting

#### **Issue #3: Continuation ID Structure in Simple Tools**

Severity: P1 - HIGH

Impact: Confusing output format, may indicate architectural issue

**Status:** UNRESOLVED (by design but questionable)

### **Root Cause:**

- Simple tools return continuation id even for single-turn operations
- Output format includes metadata that should be internal
- MCP protocol translation may be exposing internal structures

#### **Evidence:**

```
"status": "continuation_available",
"content": "Chat tool working...",
  "content_type": "text",
  "metadata": {
    "tool_name": "chat",
    "conversation_ready": true,
    "model_used": "glm-4.5-flash",
    "provider_used": "glm"
  "continuation offer": {
    "continuation_id": "62d15167-479e-4f32-9464-88c7db08b734",
    "note": "You can continue this conversation for 19 more exchanges.",
    "remaining_turns": 19
  }
}
```

### **Files Affected:**

- tools/simple/base.py (lines 400-500, response formatting)
- tools/simple/mixins/continuation mixin.py (continuation handling)
- scripts/run\_ws\_shim.py (lines 60-75, response cleaning)

### **Analysis:**

This appears to be intentional design for conversation continuity, but:

- 1. Should be optional (not forced for all responses)
- 2. Metadata should be in separate channel (not in content)
- 3. Continuation offer should only appear when explicitly requested
- 4. Format is verbose and clutters simple responses

### **Recommended Fix:**

- 1. Make continuation id optional based on request parameter
- 2. Move metadata to separate response field
- 3. Only include continuation\_offer when conversation mode is active
- 4. Simplify response format for single-turn operations

### HIGH PRIORITY ISSUES

#### Issue #4: No "wave1" Branch Exists

Severity: P1 - HIGH

Impact: Documentation references non-existent branch, cannot compare changes

Status: CONFIRMED

### **Root Cause:**

- Documentation references wave1 branch that doesn't exist in repository
- Branch comparison document (BRANCH COMPARISON wavel-to-auggie-optimization.md) references non-existent branch
- Actual branch is docs/wavel-complete-audit (different name)

### **Evidence:**

```
# From git ls-remote output:
80396b35eada5ca2d0f3703ff03a9a7319aef2de    refs/heads/docs/wave1-complete-audit
# No refs/heads/wave1 exists

# But documentation says:
"Previous Branch: docs/wave1-complete-audit"
"Current Branch: feat/auggie-mcp-optimization"
```

#### **Files Affected:**

- /home/ubuntu/Uploads/BRANCH COMPARISON wave1-to-auggie-optimization.md
- Documentation references throughout project

#### Impact:

- Cannot perform accurate branch comparison
- Unclear what "working wave1" state actually was
- May be comparing against wrong baseline

#### **Recommended Fix:**

- 1. Clarify which branch is the "working" baseline
- 2. Update documentation to reference correct branch names
- 3. Create proper branch comparison using actual branches
- 4. Document what functionality worked in baseline vs current

### **Issue #5: Timeout Configuration Chaos**

Severity: P1 - HIGH

Impact: Unpredictable behavior, difficult to tune performance

**Status:** UNRESOLVED

#### **Root Cause:**

- Multiple timeout configurations across different layers
- Auggie CLI config overrides with very long timeouts (600s+)
- No clear timeout hierarchy or documentation
- Different tools have different timeout expectations

### **Evidence:**

```
# From Daemon/mcp-config.auggie.json:
"EXAI_SHIM_RPC_TIMEOUT": "600" # 10 minutes
"EXAI_WS_CALL_TIMEOUT": "600" # 10 minutes
"KIMI_CHAT_TOOL_TIMEOUT_WEB_SECS": "900" # 15 minutes

# From config.py:
EXPERT_ANALYSIS_TIMEOUT_SECS=90 # 90 seconds

# From tools/workflows/thinkdeep.py:
THINKDEEP_EXPERT_TIMEOUT_SECS=25 # 25 seconds (default)
```

### **Timeout Hierarchy (Current - Broken):**

Tool Level: 25s (thinkdeep) / 90s (expert analysis)

Daemon Level: 600s (EXAI\_WS\_CALL\_TIMEOUT)
Shim Level: 600s (EXAI\_SHIM\_RPC\_TIMEOUT)

Provider Level: 900s (KIMI CHAT TOOL TIMEOUT WEB SECS)

**Problem:** Inner timeouts (25s, 90s) will never trigger because outer timeouts (600s, 900s) are much longer. This creates a situation where:

- 1. Tool thinks it has 25s to complete
- 2. But daemon waits 600s before timing out
- 3. User experiences 10-minute hang instead of 25s timeout
- 4. No progress updates during the wait

#### **Files Affected:**

- Daemon/mcp-config.auggie.json (Auggie CLI overrides)
- Daemon/mcp-config.augmentcode.json (VSCode Augment overrides)
- Daemon/mcp-config.claude.json (Claude Desktop overrides)
- config.py (default timeouts)
- src/daemon/ws\_server.py (daemon timeouts)
- All workflow tool files (tool-specific timeouts)

#### **Recommended Fix:**

- 1. Establish clear timeout hierarchy: Tool < Daemon < Shim < Client
- 2. Document timeout strategy in central location
- 3. Implement timeout coordination (inner timeout = 80% of outer timeout)
- 4. Add timeout warnings when approaching limits
- 5. Reduce Auggie config timeouts to reasonable values:
- Tool: 60s (simple) / 120s (workflow)

- Daemon: 180s - Shim: 240s - Client: 300s

### **Issue #6: Expert Validation Disabled**

Severity: P1 - HIGH

Impact: Workflow tools missing key feature, quality degraded

Status: KNOWN (documented in MASTER TASK LIST)

### **Root Cause:**

- Expert validation was calling analysis multiple times (duplicate calls)
- Temporarily disabled to prevent 300+ second hangs
- Bug not yet fixed, feature remains disabled

#### **Evidence:**

#### **Files Affected:**

- tools/workflow/expert analysis.py (expert validation logic)
- tools/workflow/conversation integration.py (removed stub method)
- .env (DEFAULT\_USE\_ASSISTANT\_MODEL=false)

#### Impact:

- Workflow tools don't get expert validation
- Quality of analysis reduced
- Missing key differentiator of the system

#### **Recommended Fix:**

- 1. Debug why expert analysis is called multiple times
- 2. Implement call deduplication
- 3. Add request tracking to prevent duplicate calls
- 4. Re-enable expert validation with proper safeguards
- 5. Add circuit breaker to prevent runaway calls

### MEDIUM PRIORITY ISSUES

### **Issue #7: Native Web Search Integration Unclear**

Severity: P2 - MEDIUM

Impact: Web search may not work as intended, unclear behavior

Status: PARTIALLY RESOLVED (per documentation)

#### **Root Cause:**

- Web search implementation split between GLM and Kimi
- GLM uses native web search tool (hidden from registry)
- Kimi uses \$web search builtin function
- Integration points not clearly documented in code

### **Evidence:**

### # From MASTER\_TASK\_LIST:

### ### 1.2: Web Search Integration in Chat Tool 🔽 COMPLETE

- GLM web search function is HIDDEN from tool registry (internal function only)
- AI Manager (GLM-4.5-Flash) auto-triggers web search when use\_websearch=true

#### ### 1.3: Kimi Web Search Configuration ✓ COMPLETE

- Kimi uses `\$web\_search` builtin function (correct Moonshot API format)

#### **Files Affected:**

- server.py (line 260, glm\_web\_search hidden)
- tools/simple/base.py (lines 502-508, web search auto-injection)
- src/providers/capabilities.py (lines 45-81, web search schemas)
- src/providers/orchestration/websearch adapter.py

#### **Analysis:**

According to documentation, web search is implemented and working:

- GLM: Native web search via tools schema
- Kimi: Builtin \$web search function
- Auto-injection in SimpleTool.execute()

However, code inspection shows:

- 1. Web search tool is hidden from registry (correct)
- 2. Auto-injection logic exists (correct)
- 3. But unclear if it's actually being called
- 4. No logging of web search activation
- 5. No tests verifying web search works

#### **Recommended Fix:**

- 1. Add logging when web search is activated
- 2. Add tests for web search integration
- 3. Document web search flow in architecture docs
- 4. Add metrics for web search usage
- 5. Verify web search works in both GLM and Kimi

### **Issue #8: MCP Configuration Inconsistency**

Severity: P2 - MEDIUM

Impact: Different behavior across clients, hard to maintain

**Status:** UNRESOLVED

### **Root Cause:**

- Three different MCP configurations (Auggie, Augment, Claude)
- Each has different timeout and concurrency settings
- No clear documentation of differences
- Auggie config has extreme values (600s+ timeouts)

### **Evidence:**

```
// Daemon/mcp-config.auggie.json
"EXAI_SHIM_RPC_TIMEOUT": "600"
"EXAI_WS_CALL_TIMEOUT": "600"
"EXAI_WS_SESSION_MAX_INFLIGHT": "6"

// Daemon/mcp-config.augmentcode.json
// (Different values - need to check)

// Daemon/mcp-config.claude.json
// (Different values - need to check)
```

### Files Affected:

- Daemon/mcp-config.auggie.json
- Daemon/mcp-config.augmentcode.json
- Daemon/mcp-config.claude.json

#### **Recommended Fix:**

- 1. Standardize configurations across clients
- 2. Document why differences exist (if necessary)
- 3. Create base configuration with client-specific overrides
- 4. Add validation for configuration values
- 5. Test all three configurations regularly

### **Issue #9: Bootstrap Module Complexity**

Severity: P2 - MEDIUM

Impact: Harder to maintain, potential initialization issues

**Status:** NEW ARCHITECTURE (from refactoring)

#### **Root Cause:**

- New bootstrap modules created during refactoring

- Consolidates initialization code (good)

- But adds another layer of indirection
- May have initialization order dependencies

#### **Evidence:**

```
# From scripts/run ws shim.py:
from src.bootstrap import load env, get repo root, setup logging
# Bootstrap: Setup path and load environment
load env() # Must be called first
logger = setup logging("ws shim", ...) # Then logging
```

#### Files Affected:

- src/bootstrap/\_\_init\_\_.py
- src/bootstrap/env loader.py
- src/bootstrap/logging setup.py
- All entry point scripts

#### **Analysis:**

This is actually an improvement from the previous duplicated code, but:

- 1. Initialization order matters (env → logging → everything else)
- 2. No validation of initialization success
- 3. Silent failures possible
- 4. No rollback on partial initialization

### **Recommended Fix:**

- 1. Add initialization validation
- 2. Implement proper error handling
- 3. Add initialization status tracking
- 4. Document initialization order requirements
- 5. Add tests for bootstrap module

### LOW PRIORITY ISSUES

**Issue #10: File Path Validation Too Strict** 

Severity: P3 - LOW

Impact: User experience issue, already has workaround

Status: RESOLVED (per documentation)

### **Root Cause:**

- File path validation required absolute paths
- EX ALLOW RELATIVE PATHS defaulted to false
- Fixed by changing default to true

#### **Evidence:**

```
# From COMPREHENSIVE_TOOL_TESTING_2025-10-03.md:
### Fix Applied
1. Changed default from "false" to "true" in base_tool_file_handling.py line 96
2. Added documentation to .env.example
3. Added EX_ALLOW_RELATIVE_PATHS=true to .env
```

Status: 🔽 RESOLVED

### **Issue #11: Continuation ID Expiration**

Severity: P3 - LOW

Impact: User experience issue, conversations expire

Status: BY DESIGN

### **Root Cause:**

- Conversations expire after 3 hours
- No warning before expiration
- Error message could be clearer

#### **Evidence:**

```
{
  "timestamp": 1759100315.8598037,
  "tool": "chat",
  "error": "Conversation thread 'ctx-ce818efc' was not found or has expired..."
}
```

### Files Affected:

- Conversation storage system
- Error message formatting

### **Recommended Fix:**

- 1. Add warning when conversation approaching expiration
- 2. Improve error message with recovery instructions
- 3. Consider longer expiration time (6-12 hours)
- 4. Add conversation persistence option

### **BREAKING CHANGES ANALYSIS**

### Changes from "wave1" to feat/auggie-mcp-optimization

**Note:** Cannot perform accurate comparison because "wave1" branch doesn't exist. Using docs/wave1-complete-audit as baseline.

### Major Changes (from BRANCH COMPARISON document):

- 1. Bootstrap Modules Created (NEW)
  - Consolidated initialization code

- Reduced duplication across entry points
- Impact: Positive (cleaner code)

### 2. Mixin Pattern Implementation (NEW)

- Extracted mixins from monolithic classes
- Better separation of concerns
- **Impact:** Positive (maintainability)

### 3. Auggie MCP Config Optimization (MODIFIED)

- Extended timeouts (600s+)
- Increased concurrency limits
- Impact: NEGATIVE (causes hanging issues)

### 4. Critical Bug Fixes (FIXED)

- Server crash on startup (status.py)
- Web search integration (text format handler.py)
- Legacy "zen" references removed
- Impact: Positive (stability)

#### 5. Expert Validation Disabled (CHANGED)

- Temporarily disabled due to duplicate call bug
- Impact: NEGATIVE (missing feature)

#### What Broke:

- 1. Workflow Tools: Worked in baseline, broken in current
  - Root cause: Timeout configuration changes
  - Auggie config extended timeouts too much
  - No progress heartbeat added
- 2. **Logging:** Worked in baseline, inconsistent in current
  - Root cause: Different code paths for workflow tools
  - Refactoring may have broken logging integration
- 3. Expert Validation: Worked in baseline, disabled in current
  - Root cause: Duplicate call bug discovered
  - Temporarily disabled, not yet fixed

### What Improved:

- 1. Code Organization: Better with bootstrap and mixins
- 2. **Bug Fixes:** Several critical bugs fixed
- 3. **Documentation:** Extensive documentation added (77 files)
- 4. Testing: New test infrastructure added

### API INTEGRATION ANALYSIS

### **GLM Provider (ZhipuAI/Z.ai)**

Status: WORKING

### **Configuration:**

GLM\_API\_KEY=configured GLM\_BASE\_URL=https://api.z.ai/v1 GLM\_DEFAULT\_MODEL=glm-4.6

### **Available Models:**

- glm-4.6 (flagship, 200K context)
- glm-4.5-flash (manager, fast/cheap)
- glm-4.5, glm-4.5-air, glm-4.5-x

#### Features:

- Native web search (via tools schema)
- V Streaming support
- **Tool** calling
- ✓ SDK integration (zai-sdk v0.0.4)

#### **Issues:**

- Web search tool hidden from registry (correct)
- No logging of web search activation
- Unclear if web search actually works

### **Kimi Provider (Moonshot)**

Status: WORKING

### **Configuration:**

KIMI\_API\_KEY=configured
KIMI\_BASE\_URL=https://api.moonshot.ai/v1
KIMI\_DEFAULT\_MODEL=kimi-k2-0905-preview

### **Available Models:**

- kimi-k2-0905-preview (recommended, 256K context)
- kimi-thinking-preview (deep reasoning)
- kimi-k2-turbo-preview (fast)
- Legacy: moonshot-v1-128k, moonshot-v1-32k, moonshot-v1-8k

### **Features:**

- Native web search ( \$web search builtin)
- V Streaming support
- V File upload/extract
- Advanced caching (automatic)
- OpenAl-compatible API

#### **Issues:**

- File cleanup may not be happening
- Cache management unclear
- No metrics for cache hit rate

### **Manager-First Routing**

Status: O PARTIALLY IMPLEMENTED

### Design:

```
Level 1: GLM-4.5-flash (Manager)

↓ (if complex)

Level 2: GLM-4.6

↓ (if specialized)

Level 3: Kimi
```

#### **Issues:**

- Routing logic exists but unclear if it's working
- No logging of routing decisions
- No metrics for routing effectiveness
- Manager may not be making optimal decisions

#### Files:

- src/router/service.py
- src/router/classifier.py
- src/router/unified\_router.py

### **TOOL EXECUTION FLOW ANALYSIS**

### **Simple Tools (WORKING)**

```
1. Client sends MCP request

1. 2. run_ws_shim.py receives request

1. 3. WebSocket connection to daemon (ws://127.0.0.1:8765)

1. 4. ws_server.py receives request

1. 5. server.py routes to tool

1. 6. SimpleTool.execute() runs

1. 7. Provider API call (GLM/Kimi)

1. 8. Response formatted with continuation_id

1. 9. Response sent back through WebSocket

1. 10. run_ws_shim.py cleans response

1. 11. Client receives response
```

**Timing:** 1-15 seconds (typical) **Logging:** ✓ Works correctly

Issues: Continuation\_id structure verbose

### **Workflow Tools (BROKEN)**

```
1. Client sends MCP request

1. 
2. run_ws_shim.py receives request

1. 
3. WebSocket connection to daemon

1. 
4. ws_server.py receives request

1. 
5. server.py routes to workflow tool

1. 
6. WorkflowTool.execute() runs

1. 
8. Provider API call (GLM/Kimi)

1. 
9. [HANGS HERE - No progress updates]

1. 
1. [No response or error]
```

Timing: 600+ seconds (hangs)

**Logging:** Not working

#### **Issues:**

- No progress heartbeat
- No timeout handling
- No logging during execution
- Expert validation disabled

### LOGGING SYSTEM ANALYSIS

### **Current Logging Implementation**

#### Log Files:

- .logs/toolcalls.jsonl Tool execution logs
- .logs/metrics.jsonl Performance metrics
- .logs/ws\_daemon.log Daemon logs
- .logs/ws\_shim.log Shim logs

### **Logging for Simple Tools: ✓** WORKING

```
{
  "timestamp": 1759100327.2925136,
  "tool": "chat",
  "request_id": null,
  "duration_s": 11.423,
  "result_preview": "...",
  "prompt_bullets": ["..."],
  "summary_words": 600,
  "summary_text": "..."
}
```

### Logging for Workflow Tools: NOT WORKING

- · No entries in toolcalls.jsonl when tools hang
- No progress messages captured
- · No error messages logged
- Silent failures

### **Root Cause Analysis**

### **Simple Tools:**

```
# From tools/simple/base.py
# Logging happens in execute() method
# Progress messages sent via send_progress()
# Results logged to toolcalls.jsonl
```

#### **Workflow Tools:**

```
# From tools/workflow/base.py
# Different execution path
# Progress messages may not be captured
# Logging may not be called
# Different error handling
```

Problem: Workflow tools use different code path that doesn't integrate with logging system properly.

### **CLIENT CONFIGURATION ANALYSIS**

### **VSCode Augment Configuration**

File: Daemon/mcp-config.augmentcode.json

Status: Need to inspect (not fully analyzed)

### **Expected Issues:**

- May have different timeout values
- May have different concurrency limits
- Should be tested separately

### **Auggie CLI Configuration**

File: Daemon/mcp-config.auggie.json

**Status:** • PROBLEMATIC

**Configuration:** 

```
{
  "EXAI_SHIM_RPC_TIMEOUT": "600",
  "EXAI_WS_CALL_TIMEOUT": "600",
  "EXAI_WS_SESSION_MAX_INFLIGHT": "6",
  "EXAI_WS_GLOBAL_MAX_INFLIGHT": "16",
  "KIMI_CHAT_TOOL_TIMEOUT_WEB_SECS": "900",
  "EX_SESSION_SCOPE_STRICT": "false",
  "EX_SESSION_SCOPE_ALLOW_CROSS_SESSION": "true"
}
```

#### Issues:

- 1. Timeouts Too Long: 600s (10 min) causes hanging perception
- 2. Concurrency Reduced: From 12 to 6 (may be intentional)
- 3. Session Scope Relaxed: May cause cross-contamination

#### Rationale (from documentation):

- Extended timeouts for "max thinking mode"
- Support 30-60 minute autonomous sessions
- Better workflow continuity

**Problem:** While rationale makes sense, implementation causes:

- User perceives system as hanging
- No progress updates during long waits
- Timeout hierarchy broken
- No graceful degradation

### **Claude Desktop Configuration**

File: Daemon/mcp-config.claude.json

Status: Need to inspect (not fully analyzed)

### **Expected Issues:**

- May have different timeout values
- May have different concurrency limits
- Should be tested separately

# **CODE QUALITY ASSESSMENT**

## **Architectural Strengths**

- 1. Modular Design: Good separation between providers, tools, routing
- 2. Mixin Pattern: Clean extraction of concerns into mixins
- 3. Bootstrap Modules: Consolidated initialization code
- 4. Provider Abstraction: Clean interface for GLM and Kimi
- 5. MCP Protocol: Proper implementation of MCP standard

### **Architectural Weaknesses**

- 1. **Timeout Management:** Scattered across multiple files, no coordination
- 2. Logging Inconsistency: Different paths for simple vs workflow tools

- 3. Error Handling: Silent failures, no graceful degradation
- 4. Progress Reporting: Missing for long-running operations
- 5. Configuration Complexity: Too many environment variables, unclear hierarchy

### **Code Smells**

- 1. Magic Numbers: Timeout values hardcoded in multiple places
- 2. Duplicate Logic: Timeout handling duplicated across tools
- 3. Silent Failures: Errors not propagated properly
- 4. Missing Validation: Configuration values not validated
- 5. **Inconsistent Patterns:** Simple vs workflow tools use different patterns

### **Design Issues**

- 1. **Tight Coupling:** Timeout configuration tightly coupled to implementation
- 2. Missing Abstraction: No timeout manager or coordinator
- 3. No Circuit Breaker: Repeated failures not handled
- 4. No Health Checks: Long operations don't report health
- 5. No Metrics: Can't measure system performance

### PRIORITY RANKING

### PO - CRITICAL (Must Fix Immediately)

- 1. Workflow Tools Hanging (Issue #1)
  - Impact: System unusable for complex tasks
  - Effort: Medium (2-3 days)
  - Fix: Implement timeout hierarchy + progress heartbeat
- 2. Logging Not Working (Issue #2)
  - Impact: Cannot debug issues
  - Effort: Medium (2-3 days)
  - Fix: Unify logging infrastructure

### P1 - HIGH (Fix Soon)

- 1. Timeout Configuration Chaos (Issue #5)
  - Impact: Unpredictable behavior
  - Effort: Medium (2-3 days)
  - Fix: Standardize timeout hierarchy
- 2. Expert Validation Disabled (Issue #6)
  - Impact: Missing key feature
  - Effort: High (3-5 days)
  - Fix: Debug duplicate call issue

#### 3. Branch Comparison Issue (Issue #4)

- Impact: Cannot verify changes
- Effort: Low (1 day)
- Fix: Clarify baseline branch

### P2 - MEDIUM (Fix When Possible)

### 1. Continuation ID Structure (Issue #3)

- Impact: Confusing outputEffort: Low (1-2 days)
- Fix: Make optional, simplify format

#### 2. Web Search Integration (Issue #7)

- Impact: Feature may not work
- Effort: Medium (2-3 days)
- Fix: Add logging, tests, verification

### 3. MCP Config Inconsistency (Issue #8)

- Impact: Different behavior across clients
- Effort: Medium (2-3 days)
- Fix: Standardize configurations

### P3 - LOW (Nice to Have)

### 1. Bootstrap Module Complexity (Issue #9)

- Impact: Maintenance burden
- Effort: Low (1-2 days)
- Fix: Add validation, tests

### 2. Continuation Expiration (Issue #11)

- Impact: User experience
- Effort: Low (1 day)
- Fix: Better warnings, longer expiration

# **RECOMMENDED FIX SEQUENCE**

### Phase 1: Critical Fixes (Week 1)

Goal: Make workflow tools functional

### 1. Day 1-2: Implement Timeout Hierarchy

- Define clear timeout hierarchy
- Implement timeout coordination
- Add timeout warnings
- Test with workflow tools

### 2. Day 3-4: Add Progress Heartbeat

- Implement progress reporting for long operations
- Add heartbeat every 5-8 seconds
- Test with workflow tools
- Verify user experience

### 3. Day 5: Fix Logging

- Unify logging infrastructure
- Ensure workflow tools log correctly
- Add structured logging
- Test logging for all tools

### **Phase 2: High Priority Fixes (Week 2)**

Goal: Restore full functionality

### 1. Day 6-8: Fix Expert Validation

- Debug duplicate call issue
- Implement call deduplication
- Re-enable expert validation
- Test with workflow tools

### 2. Day 9-10: Standardize Configurations

- Create base configuration
- Standardize timeout values
- Document configuration hierarchy
- Test all three clients

### Phase 3: Medium Priority Fixes (Week 3)

Goal: Improve reliability and usability

### 1. Day 11-12: Simplify Continuation System

- Make continuation id optional
- Simplify response format
- Move metadata to separate field
- Test with all tools

#### 2. Day 13-14: Verify Web Search

- Add web search logging
- Create web search tests
- Verify GLM and Kimi integration
- Document web search flow

### 3. Day 15: Documentation Update

- Update all documentation
- Document timeout hierarchy
- Document configuration options
- Create troubleshooting guide

### **TESTING RECOMMENDATIONS**

### **Unit Tests Needed**

### 1. Timeout Management

- Test timeout hierarchy
- Test timeout coordination
- Test timeout warnings
- Test graceful timeout

### 2. Logging System

- Test simple tool logging
- Test workflow tool logging

- Test structured logging
- Test log rotation

### 3. Continuation System

- Test continuation creation
- Test continuation retrieval
- Test continuation expiration
- Test cross-session continuations

### **Integration Tests Needed**

#### 1. Workflow Tools

- Test analyze tool end-to-end
- Test thinkdeep tool end-to-end
- Test debug tool end-to-end
- Test codereview tool end-to-end

#### 2. Web Search

- Test GLM web search
- Test Kimi web search
- Test web search auto-injection
- Test web search results

### 3. Expert Validation

- Test expert validation flow
- Test duplicate call prevention
- Test expert timeout handling
- Test expert result integration

### **Performance Tests Needed**

### 1. Timeout Behavior

- Measure actual timeout values
- Test timeout cascade
- Test progress heartbeat frequency
- Test graceful degradation

### 2. Concurrency

- Test session concurrency limits
- Test global concurrency limits
- Test provider concurrency limits
- Test concurrent workflow tools

### 3. Logging Performance

- Measure logging overhead
- Test log file size growth
- Test log rotation
- Test log parsing performance

### CONCLUSION

The EX-AI MCP Server on the feat/auggie-mcp-optimization branch has **critical architectural issues** that prevent complex workflow tools from functioning correctly. While simple tools work, the system is **not production-ready** for complex use cases.

### **Key Takeaways**

- 1. Simple Tools Work: Chat and other simple tools function correctly
- 2. Workflow Tools Broken: Analyze, thinkdeep, debug, codereview hang without timeout
- 3. Logging Inconsistent: Works for simple tools, broken for workflow tools
- 4. Timeout Chaos: Multiple conflicting timeout configurations
- 5. Expert Validation Disabled: Key feature temporarily disabled due to bug
- 6. Documentation Mismatch: References non-existent "wave1" branch

### **Root Causes**

- 1. Timeout Hierarchy Broken: Inner timeouts never trigger due to outer timeouts being too long
- 2. No Progress Heartbeat: Long operations don't report progress
- 3. Logging Path Divergence: Workflow tools use different code path without logging
- 4. Configuration Overrides: Auggie config extends timeouts too much (600s+)
- 5. Missing Error Handling: Silent failures instead of graceful degradation

### **Immediate Actions Required**

- 1. Fix Timeout Hierarchy: Implement proper timeout coordination
- 2. Add Progress Heartbeat: Report progress every 5-8 seconds
- 3. Fix Logging: Unify logging infrastructure for all tools
- 4. **Reduce Timeouts:** Change Auggie config to reasonable values (90-180s)
- 5. **Test Workflow Tools:** Verify all workflow tools work correctly

### **Long-Term Improvements**

- 1. Standardize Configurations: Create base config with client overrides
- 2. Fix Expert Validation: Debug and re-enable expert validation
- 3. Add Monitoring: Implement health checks and metrics
- 4. Improve Documentation: Update all docs with accurate information
- 5. Add Tests: Create comprehensive test suite

**Report Generated:** 2025-10-04 **Analyst:** Abacus.Al Deep Agent

**Repository:** https://github.com/Zazzles2908/EX-AI-MCP-Server

**Branch:** feat/auggie-mcp-optimization

**Commit:** e27cf6f (docs: add branch comparison documentation)